

# Ian Meyer

Full remote, US-based

(917) 803-8928

[ianmmeyer@gmail.com](mailto:ianmmeyer@gmail.com)

## **ReadySet Technologies - Remote**

### **Staff Site Reliability Engineer - May 2022 to July 2023**

---

- AWS Cost Resource Management – Reduce monthly spending by ~\$15K across multiple accounts and environments within a 2 week period. Started as a manual process eventually moving to an automated recurring process using the `aws-nuke` utility to prevent costly orphaned resources.
- Implement observability best practices – Updates to core to export metrics on a standard path. Create Grafana dashboard for end-user monitoring. Identify SLIs within exposed metrics to build a SLO profile
- Rewrite Helm chart for more flexible deployments in Kubernetes
- Provide project management and leadership to a vertical team for sixth months. Focus on leveling up junior platform engineers, and providing better timelines for project delivery.
- Perform HR duties related to onboarding/off-boarding of engineers
- Onboard contracted IT team to provide helpdesk management resources, as well as take on most onboard/off-board duties

## **Workato - Remote**

### **Principal Site Reliability Engineer - Nov 2021 to May 2022**

---

- Onboarding was the only thing accomplished. Workato, being a largely Russian-employed org, the Russia/Ukraine war significantly impacted operations.

## **DivvyDose/Optum RX - Remote**

### **Principal Release Engineer - Jan 2021 to Nov 2021**

---

- Evangelizing best practices for feature launches, release pipelines, and reliability within the SDLC through leading regular on-boarding sessions, holding regular office hours with the Platform team (my team) as the host
- Provided architectural guidance to engineers on decomposing their monolith into lambdas or micro-services
- Migrated legacy build jobs on Jenkins over to CircleCI and Github Actions
- Built an internal tool to track service deployments using Golang, shortly after which...
- Migrated all deployment tracking services to DataDog CI platform
- Implemented SonarCloud to enable developers to commit better, and more secure code to reduce security risks through outdated/at-risk dependencies or other code violations (not disabling TLS1 etc)
- Provided AWS accounts (as a service) to product engineering teams, with established SCP to ensure HIPAA compliance
- Assisted data engineering team in evaluating tooling/vendors for analyst/data engineer workflows

## **build.com / Ferguson - Remote**

### **Senior Site Reliability Engineer - Jan 2020 to Dec 2020**

---

- Replace legacy monitoring infrastructure with new platform based on Prometheus/Alertmanager
- Work with Ferguson leadership on how best to integrate the build.com platform
- Build tooling for automating deployment of monitoring platform components
- Automated RPM based deployment by datacenter using Jenkins/[harness.io](https://harness.io)

## **MediaMath – Remote/New York, NY**

### **Staff Site Reliability Engineer - Jan 2015 to Dec 2019**

---

- Built in-house tool using Go/BoltDB to store and track SLO performance over time
- Created an in-house, self-service monitoring platform based on Prometheus and Alertmanager used within the engineering organization.

- Re-architected deployment system of APIs to enable zero-downtime deploys using runit and perl's starman web-server with canaries for individual services
- Standardized all API systems with Chef using custom cookbooks specific to our infrastructure
- Mentor SREs on best practices for balancing risk-mitigation and feature deployment. (feature flags, canaries, dark launches)
- Implemented Prometheus+Alertmanager for monitoring/alerting of critical API system resources and metrics (histograms to track response latency, linear regressions against local storage utilization to alert on filling a disk within 6h, among others)
- Upgraded API infrastructure to maintain SOC-2 compliance efforts
- Help engineers transition from monolithic perl-based applications to modern microservices stack (golang, scala) via deep-dives into critical business dependencies within the underlying data model

#### **Percolate Industries – Remote/New York, NY**

##### **Senior Infrastructure Engineer - May 2014 to Dec. 2014**

---

- Lead postmortems for production incidents, as well as assuming the incident commander role during live incidents.
- Deployed Percolate internal app (4+ times daily)
- Created utility in Go to monitor remote Github repositories for branches older than D duration.
- Discovered, and fixed bug in deploy reducing time to deploy to ~10 minutes, from ~30 minutes.
- Assisted other engineers with profiling and monitoring their services.

#### **Knewton, Inc - Remote/New York, NY**

##### **Director of Systems Engineering - Apr. 2013 to Apr. 2014**

---

- Led team of seven systems engineers (half remote, half local to NYC) in building out an internal platform for managing and deploying apps.
- Create unified team vision and objectives and communicate to rest of organization.
- Lead postmortems for production incidents, and also communication during incidents.
- Create interview process for new systems engineers
- Model AWS costs 6/12/18+ months in the future
- Perform quarterly team performance reviews

#### **Admeld/Google, Inc - Remote/New York, NY**

##### **Manager, Site Reliability Engineer - Nov. 2010 to Feb. 2013**

---

- Led Admeld's technical integration for production ad-serving environments with Google engineers
- Handled migration of vendor relationships over to Google
- On-board Google engineers for working with the Admeld production infrastructure.
- Integrated Google processes for revision control, service monitoring, alert escalation, and infrastructure management with legacy Admeld infrastructure.
- Modeled capacity for scaling up and down the legacy Admeld infrastructure.
- Leveraged Chef as Admeld's configuration management choice.
- Full implementation of cookbooks to create and manage Admeld infrastructure for ~50 analytics/reporting services written in Java and 3 ad-serving components written in C++.
- Created standardized hardware/software platform for networking gear, physical infrastructure, operating system, and operating system configuration.
- Led a case study of Admeld for Opscode/Chef, highlighting the value added to our operations org.
- Built out data-warehouse datacenter in NYC. / Oversaw disaster remediation during Hurricane Sandy.

#### **Etsy, Inc – New York, NY**

##### **Senior Systems Administrator - Apr 2009 to June 2010**

---

- Implemented LDAP across production and corporate networks to facilitate a unified login across servers/applications.
- Installed Cobbler to manage system provisioning and assist in OS consolidation to a RedHat based infrastructure.
- Set up Graphite to capture business level metrics such as revenue, listings sold, new users, etc.
- Decoupled our DNS, DHCP, kickstart provisioning system into standalone services.
- Designed, and specified configuration of HP DL785s used for migration of our master database system.
- Streamlined our syslog infrastructure to centralized log aggregation with syslog-ng and Splunk.
- Reduced the number of database calls using the PyPgSQL python module by hardcoding the PostgreSQL version in the C code.
- Participated in a small team with technical executives to address the top issues facing our organization. Some things to come out of this group was the shift to continuous deployment, SOA, and a more collaborative team-based collaboration.
- Wrote custom gmetric scripts to inject application level metrics into Ganglia that measure number of requests per second, and the average response time for a 15 second period for apache.
- Lead the migration from Ubuntu to Fedora Core 10 for our front end-serving platform.
- Installed, configured, deployed a development, staging and production environment for Chef across 100+ servers.

#### **RightMedia/Yahoo – New York, NY**

##### **Production Engineer III – Aug 2006 to April 2009**

---

- Assisted in design and implemented a system to support processing of over 250,000 requests per second through the Right Media Adexchange, across a global environment Expanded from 300 production servers to over 12,000 production servers, spread across 6 data centers throughout the world in 24 months.
- Designed kickstart system coupled with CFEngine to provide autonomous installation and postinstallation configuring of server and application. Architected and implemented a 400+ user LDAP system with global redundancy utilizing a master/ slave environment replicating changes to 12 slaves worldwide within minutes.
- Performed daily backups of backend data systems using Bacula.
- Implemented DNS views with BIND to simplify configuration of servers across data centers.
- Monitored all backend data and front end production systems with Nagios, implementing escalation paths within specific intervals for oncall rotations in the US and India.
- Using Foundry load-balancers and UltraDNS, managed global redundancy to prevent outages in specific regions.
- Diagnosed performance issues alongside adserver engineers to decrease memory footprint negating the need to move to 16GB memory on hosts as well as decrease latency by 50ms overall.
- Migrated all production servers from CentOS to RHEL 4.6 over the course of 1 month using standard Yahoo installation and configuration management tools. Implemented simple application installation by converting from proprietary Yahoo package format to RPM.
- Brought all production sites up to > 99.98% availability for external monitoring with Gomez.
- Implemented a new cluster naming scheme to better facilitate management and deployment of new versions of the serving application.
- Trained a team of Service Engineers in India to manage serving environment in US off hours.
- Used RRD to graph performance metrics for all servers in aggregate, segmented by data center.